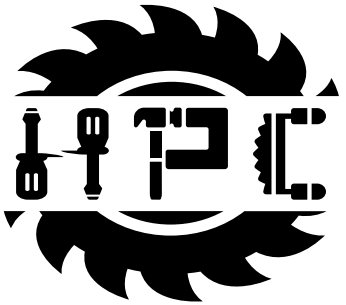

CARPENTRY

HPC Carpentry

A scalable, peer-reviewed training program to democratize HPC access

Lightning Talk

SCA 2024 Supercomputing Asia

Feb. 19-22, 2024

Our Team:

- Andrew Reid (he/him), US NIST
- Annajiat Alim Rasel (he/him), BRAC U
- Alan O’Cais (he/him), CECAM
- Trevor Keller (they/them), US NIST
- Wirawan Purwanto (he/him), ODU

HPC Carpentry

A scalable, peer-reviewed training program to democratize HPC access

Outline

This is necessarily a brief introduction to our effort

- General overview of the Carpentries
- How HPC Carpentry fits into and builds upon the Carpentries
- Brief history of the HPC Carpentry effort
- Our Strategic Plan
- Challenges

The Carpentries

Motivated by a common problem in research practice – increasingly foundational computational research work was done by a variety of practitioners, often short-timers, with a variety of skill sets, and a variety of practices.

Founded by Greg Wilson starting in the late 1990s, with various structural changes in response to feedback over the years – current umbrella organization (2018) encompasses several sub-organizations – software carpentry, library carpentry, and data carpentry.

Central idea: The skills necessary to demonstrate the value of better practices can be taught in a workshop setting working hands-on with the actual systems – learners build "muscle memory" of having done key steps in the process, and come away with the power to move themselves forward.

- Pedagogically sound – relies on peer-reviewed educational literature
  - Formal instructor training and instructor trainer training
- Open-source – uses instructor and learner feedback to modify lessons
- Inclusive – anyone can contribute feedback, pull-requests

HPC Carpentry

Common challenges with HPC environment:

- Researchers with HPC requirements often do not have adequate training in operating HPC systems, and struggle to effectively scale up/out their operations even when the hardware resources are available.

- HPC facility operators are frustrated by non-knowledgeable users making inappropriate use of shared resources, disrupting other users, and requiring significant hand-holding before they can become effective.

⇨ Candidate solution: Use Carpentries techniques to build the bridge for HPC users.

Central idea: The skills necessary to demonstrate the value of better practices can be taught in a workshop setting working hands-on with the actual systems – learners build “muscle memory” of having done it, and come away with the power to move themselves forward.

CARPENTRY

HPC Carpentry

Some Important Contributions

- ComputeCanada (now Digital Research Alliance of Canada) contributions
- Peter Steinbach, “HPC In A Day”
- SC17 BoF
  - Andy Turner, Christina Koch, Tracy Teal, Bob Freeman, Chris Bording
- CarpentryCon 2018 discussions
- SC18 BoF
  - Andy Turner, Christina Koch, Peter Steinbach, Alan O’Cais, Jeffrey Stafford, John Simpson, Daniel Smith, Bob Freeman
- Well-attended informal session at SC19
- CarpentryCon 2020@Home
  - Trevor Keller, Christina Koch, and others
- SC21 BoF – high-value feedback from HPC operators
- Present in the Carpentries Incubator, Winter 2021
- Begin working towards Lesson Program Incubation
- CarpentryCon 2022
  - Lightning talk, sprint, breakout session
- June 2023 – formal acceptance into Lesson Program Incubation

HPC Carpentry

Strategic Plan

- Build two two-day workshop tracks
  - UNIX Shell*, HPC Intro, and HPC Workflow† lessons, with Admahl's Law application, for HPC users
  - UNIX Shell*, HPC Intro, and HPC Development†† lessons, for coders to learn parallel frameworks
- Identify learner communities and hold workshops
- Reach out to other organizations focused on HPC education
- Continue to curate and host additional material across a spectrum of use cases

Our Current Lessons

- HPC Intro (Queuing system basics) (Beta, In Carpentries Incubator)
- HPC Workflow (Alpha, teachable but preliminary)
- HPC Parallel Novice (Parallelization using Python)
- HPC Shell (deprecated in favor of the Carpentries UNIX Shell lesson)
- HPC Chapel†

- Notes: * Software Carpentry. † Under development. †† Planned.

Our principal audience is novice HPC users. Where possible and practical, we are also interested in serving less-novice HPC users, as well as HPC facility operators. The more community participates, the better we can do together!

HPC Carpentry

Challenges

- Finding learners
  - We think we address an important problem, but not everyone agrees
  - Many HPC centers have existing training they like
    - But it may not be suitable for beginners?
- Ensuring workshops have clusters to teach on!
  - We have a Jetstream2 allocation we are using for this
    - But it’s time-limited, expires in a year
  - We have Terraform/Magic Castle tools to stand up clusters
    - But orgs might have to pay for the cloud to do it on
- Making lessons focused
  - We don’t want too many “black boxes”, so it’s tempting to dive in to conversations about MPI, Dask, etc.
  - But the focus of the lesson is cluster ops, not code!

We have a motivated and skilled team, and feel that these are all fixable with continued outreach and connection in to the broader HPC Education world!

HPC CARPENTRY

HPC Carpentry

Getting Involved

The GitHub Project Page:
`github.com/hpc-carpentry`

The Slack:
`#hpc-carpentry` on **`carpentries.slack.com`**
`( slack-invite.carpentries.org )`

The Topicbox E-mail List:
`carpentries.topicbox.com/groups/discuss-hpc`

The main website:
`hpc-carpentry.org`

Alan O'Cais is supported by the Digital Europe Programme under Grant No. 101100604 (BioNT)
HPC Carpentry logo by Mumtahina Sistry, BRAC U (Jan 2021)